

# SpinLock Usage in Loop

## List of Lock in Loop

```
./core/kernel/kthread.c:107:partition_t *CreatePartition(struct xmcPartition *cfg) {

./core/objects/ttnocports.c:378:xm_s32_t __VBOOT SetupComm(void) {
./core/objects/commports.c:1076:xm_s32_t __VBOOT SetupComm(void) {

./core/include/kthread.h:170:static inline void SetPartitionHwIrqPending(partition_t *p, xm_s32_t
irq) {
./core/include/kthread.h:205:static inline int ArePartitionExtIrqPendingSet(partition_t *p,
xm_s32_t irq) {
./core/include/kthread.h:239:static inline void SetPartitionExtIrqPending(partition_t *p, xm_s32_t
irq) {

./core/include/sched.h:158:static inline void SUSPEND_PARTITION(xmId_t id) {
./core/include/sched.h:170:static inline void RESUME_PARTITION(xmId_t id) {
./core/include/sched.h:198:static inline void HALT_PARTITION(xmId_t id) {

./core/kernel/mmu/physmm.c:189:void PmmResetPartition(partition_t *p) {

./core/kernel/mmu/vmmap.c:127:xmAddress_t SetupPageTable(partition_t *p, xmAddress_t
pgTb, xmSize_t size) {
./core/kernel/mmu/hypercalls.c:161:static void SetPtd(xmAddress_t pAddr, struct physPage
*pagePtd, xm_u32_t type) {
./core/kernel/mmu/hypercalls.c:188:static void SetPte(xmAddress_t pAddr, struct physPage
*pagePte, xm_u32_t type) {
./core/kernel/mmu/hypercalls.c:123:static void UnsetPtd(xmAddress_t pAddr, struct physPage
*pagePtd, xm_u32_t type) {
./core/kernel/mmu/hypercalls.c:142:static void UnsetPte(xmAddress_t pAddr, struct physPage
*pagePte, xm_u32_t type) {
```

## SpinLock Implementation

### SpinLock

```
static inline void __ArchSpinLock(archSpinLock_t *lock) {
#ifdef CONFIG_SMP
    __asm__ __volatile__ ("n1:\n\t" \
        "ldstuba [%0] 1, %%g2\n\t" /* ASI_LEON23_DCACHE_MISS */ \
        "orcc %%g2, 0x0, %%g0\n\t" \
        "bne,a2f\n\t" \
        "ldub [%0], %%g2\n\t" \
```

```

        ".subsection 2\n" \
        "2:\n\t" \
        "orcc %%g2, 0x0, %%g0\n\t" \
        "bne,a 2b\n\t" \
        "ldub [%0], %%g2\n\t" \
        "b,a 1b\n\t" \
        "nop\n\t" \
        ".previous\n" \
        : /* no outputs */ \
        : "r" (&lock->lock) \
        : "g2", "memory", "cc");
#endif
}

```

## SpinUnlock

```

static inline void __ArchSpinUnlock(archSpinLock_t *lock) {
#ifdef CONFIG_SMP
    __asm__ __volatile__ ("stb %%g0, [%0]" : : "r" (&lock->lock) : "memory");
#endif
}

```

## Who Uses SpinLock

\$ grep -lr "SpinLock" .

./core/drivers/leon\_uart.c # → checked; not likely to have “Lock in Loop”

./core/kernel/mmu/physmm.c # → **Case 3**

./core/kernel/arch/leon\_timers.c # → checked; all commented

./core/objects/console.c # → checked; less likely to have “Lock in Loop”

./core/objects/ttnocports.c # → **Case 1**

./core/objects/commports.c # → **Case 1**

./core/klibc/stdio.c # → checked; not likely to have

./core/include/spinlock.h

./core/include/list.h # → **Case 1**

./core/include/arch/spinlock.h

./core/include/physmm.h # → **Case 3**

./core/include/kthread.h # → **Case 2**

./core/include/logstream.h # → **Case 4**

# Case 1: List

```
struct dynList {  
    struct dynListNode *head;  
    xm_s32_t noElem;  
    spinLock_t lock;  
};
```

```
static inline void DynListInit(struct dynList *l) {  
    l->lock=SPINLOCK_INIT;  
    SpinLock(&l->lock);  
    l->noElem=0;  
    l->head=0;  
    SpinUnlock(&l->lock);  
}
```

```
static inline xm_s32_t DynListInsertHead(struct dynList *l, struct dynListNode *e) {  
    SomethingHere()  
    SpinLock(&l->lock);  
    SomethingHere()  
    SpinUnlock(&l->lock);  
    return 0;  
}
```

Similar SpinLock & SpinUnlock for other functions in core/include/spinlock.h

And also for syntax sugar

```
DYNLIST_FOR_EACH_ELEMENT_BEGIN  
DYNLIST_FOR_EACH_ELEMENT_END
```

## Who uses DynList

./core/kernel/ktimer.c # → checked; No “Lock in Loop”

./core/kernel/mmu/physmm.c # → checked; functions are less likely to be called in loop

./core/kernel/kthread.c # → **function CreatePartition()**

./core/objects/ttnocports.c # → **function SetupComm()**

./core/objects/commpports.c # → **function SetupComm()**

## SetupComm

```
/* create the channels */
```

```
for (e=0; e<xmcTab.noCommChannels; e++) {  
    switch(xmcCommChannelTab[e].type) {  
        case XM_QUEUING_CHANNEL:
```

```

    GET_MEMZ(channelTab[e].q.msgPool, sizeof(struct
msg)*xmcCommChannelTab[e].q.maxNoMsgs);
    DynListInit(&channelTab[e].q.freeMsgs);
    DynListInit(&channelTab[e].q.recvMsgs);
    for (i=0; i<xmcCommChannelTab[e].q.maxNoMsgs; i++) {
        GET_MEMZ(channelTab[e].q.msgPool[i].buffer,
xmcCommChannelTab[e].q.maxLength);
        if(DynListInsertHead(&channelTab[e].q.freeMsgs,
&channelTab[e].q.msgPool[i].listNode)) {
            SystemPanic();
        }
    }
}
}
}

```

**Channel** is a union structure, used for representing Sampling Channel, TTNOC Channel or Queuing Channel.

Only Queuing Channel needs a DynList for storing Msg.

GET\_MEMZ located maxNoMsgs of msg to channel message pool.

And feed all the dummy msg into freeMsgs DynList

This is a “Lock in Loop” condition.

## Case 2: *kThread\_t* Flags & *IrqPend*

### kthread.c & kthead.h

#### kthread.c

There is one function called

**static inline** *kThread\_t* \*AllocKThread(*xmld\_t* id);

which used **DynListInit()**; and this function is called from **CreatePartition()** function's loop.

```

for (i=0; i<cfg->noVCpus; i++) {
    SomethingHere();
    p->kThread[i]=k=AllocKThread(PART_VCPU_ID2KID(cfg->id, i));
    SomethingHere();
}

```

#### kthead.h

core/include/kthread.h

First of all, *kThread* contains *spinLock\_t* for sure.

Secondly,

// Write to Flags

**SetKThreadFlags**

**ClearKThreadFlags**

// Read from Flags

**AreKThreadFlagsSet**

These 3 functions contain similar format of using SpinLock and SpinUnlock. for example:

```
static inline void SetKThreadFlags(kThread_t *k, xm_u32_t f) {  
    SpinLock(&k->ctrl.lock);  
    k->ctrl.flags |= f;  
    if (k->ctrl.g && k->ctrl.g->partCtrlTab)  
        k->ctrl.g->partCtrlTab->flags |= f;  
    SpinUnlock(&k->ctrl.lock);  
}
```

The reason we need SpinLock when setting kThread Flags is that XtratuM is using Flags to check current thread running status, for example *KTHREAD\_HALTED\_F*, *KTHREAD\_TRAP\_PENDING\_F*, etc.

Not even flags, HwIrq Pending and ExtIrq Pending signals also need SpinLock to make sure the IRQ is passed to partitions correctly.

So the functions:

**SetPartitionHwIrqPending** # → has “Lock in Loop”

SetExtIrqPending

**ArePartitionExtIrqPendingSet** # → has “Lock in Loop”

AreExtIrqPendingSet

**SetPartitionExtIrqPending** # → has “Lock in Loop”

The “Lock in Loop” conditions happens when it is setting/checking partitions’ IRQ. It is needed to have a for loop like following.

```
for (e = 0; e < p->cfg->noVCpus; e++) {  
    FlagOperationHere();  
}
```

The flag operations mentioned above are used in other for loop as well, such as core/include/sched.h:1158 SUSPEND\_PARTITION; 1170 RESUME\_PARTITION  
Because of noVCpus → nokThread, it needs a for loop to change each kThread.

## Case 3: Physical Memory Manager

### core/kernel/mmu/physmm.c

#### void PmmResetPartition(partition\_t \*p)

```
for (e=0; e<p->cfg->noPhysicalMemoryAreas; e++) {
    SomethingHere();
    if (memRegion->flags&XMC_REG_FLAG_PGTab)
        for (addr=memArea->startAddr; addr<memArea->startAddr+memArea->size;
            addr+=PAGE_SIZE) {
            SomethingHere();
            SpinLock(&page->lock);
            page->type=PPAG_STD;
            page->counter=0;
            SpinUnlock(&page->lock);
            SomethingHere();
        }
}
```

Each page (struct phyPage) contains one spinlock\_t

```
struct physPage {
    struct dynListNode listNode;
#ifdef CONFIG_ARCH_MMU_BYPASS
    xmAddress_t vAddr;
#endif
    xm_u32_t mapped:1, unlocked:1, type:3, counter:27;
    spinLock_t lock;
};
```

PmmResetPartition is called from ResetPartition, and before which, all the VCpus are halted

### core/include/phymm.h

#### PPagIncCounter & PPagDecCounter

```
SpinLock(&page->lock);
cnt=page->counter;
page->counter++;
SpinUnlock(&page->lock);
```

Function PPagIncCounter is called from a loop in core/kernel/mmu/vmmap.c: SetupPageTable() and loop in core/kernel/mmu/hypercalls.c: SetPtd() SetPte()

Function PPagDecCounter is called from a loop in core/kernel/mmu/hypercalls.c: UnsetPtd() UnsetPte()

## Case 4: LogStream

### LogStreamInsert

core/include/logstream.h

```
static inline xm_s32_t LogStreamInsert(struct logStream *IS, void *log) {  
    SomethingHere();  
    SpinLock(&IS->lock);  
    SomethingHere();  
    SpinUnlock(&IS->lock);  
}
```

Though other LogStream operations are in this format too, this insert function is more likely to be used in a loop.

It is found that this function is called in a loop from function **WriteTrace()**, from core/objects/trace.c

The reason of using SpinLock may be because the health monitor is shared by different partitions. Hence, the operations of log need well-controlled.

## Summary

XtratuM does have several “Lock in Loop” conditions. Luckily, it is found that SpinLock and SpinUnlock are always used properly. There is always a SpinUnlock() before return or at the corresponding position of SpinLock().